

Email Authorship Identification

(on Enron dataset)

Automated Learning and Data Analysis

Group P09

Mohammad Ali Khaksar – mkhaksa

Sumitosh Pal – spal3


Juan Jose Rivera – jrivera5

Preya Shabrina – pshabri

Project Idea

Email Corpus

After further thought, it seems to me that in light of our fear of litigation w/ American coal, we should keep the documents. To further insulate the Coal Group and you from any claim that Enron misused the information, I suggest that you transfer the information to me and I will hold it for safekeeping.



- Our project is about identifying author of an email, using the stylometric characteristics of the email corpus.
- And we have also compared different classification algorithms.

Related Works

- When was the first attempt taken to address the problem?
- The oldest research we found was in 2003
 - **Nizamani and Memon - cluster based classification technique**
 - **Zheng et al. applied inductive learning algorithms with message features including style markers**
 - **Chen et al. proposed a frequent pattern and machine learning based model**
- May be the first attempt was taken long before 2003

Then why are we doing this again?

Motivation



Email Phishing scams cost American business 100 Billion dollars a year!!!

Motivation



Google and Facebook got duped out of **100 million dollars** through an email phishing scheme when a hacker impersonated a computer-parts vendor!!!

Motivation

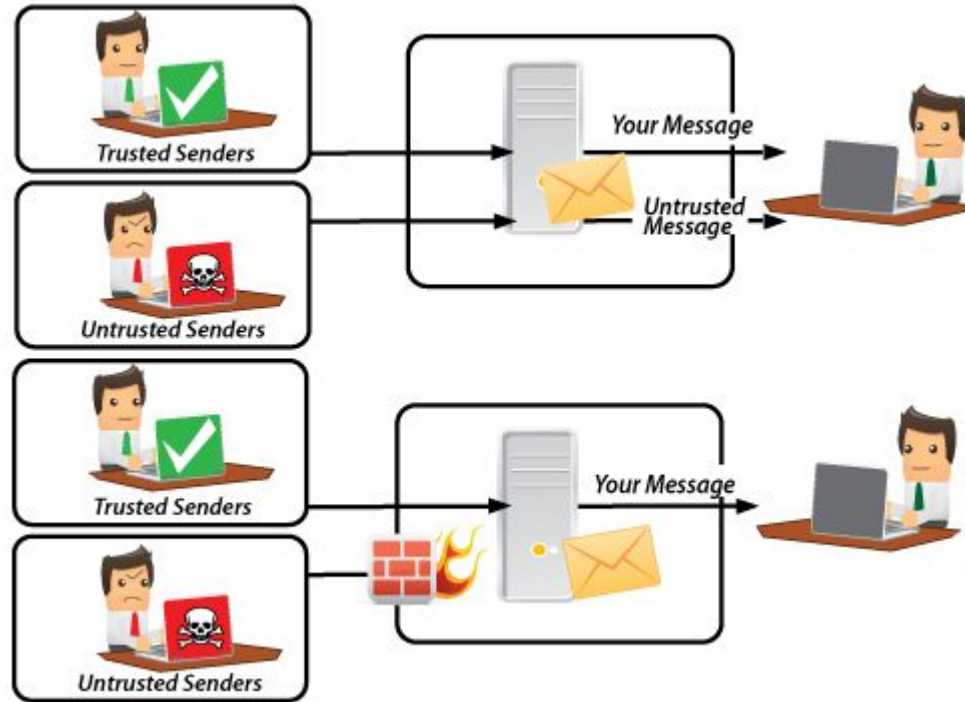


According to the FBI, criminals made off with at least **676 million** thanks to so-called business email compromise campaigns, which are attacks designed to trick company executives or accounting departments into sending money to fake vendors.

Motivation

All these incidents happened in 2018!!!

Motivation



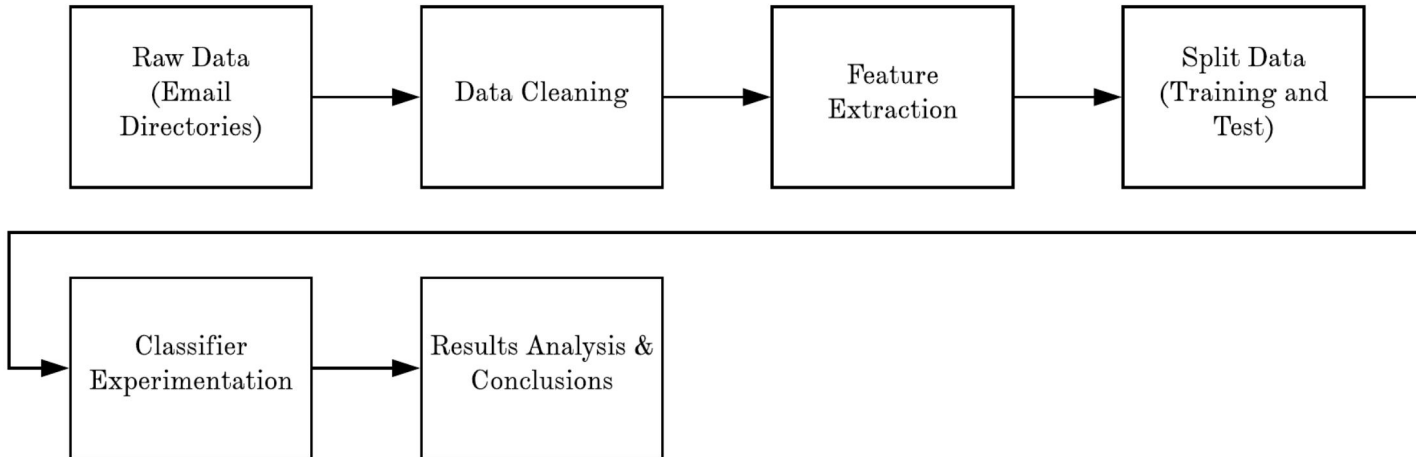
- Email authorship identification is important to categorize emails from trusted and untrusted sources
- It's also important for identity validation
- It's hard but selecting an approach is even harder

What's new?

We tried to answer the following research questions:

- Will Delta metrics Algorithms together with feature extraction mechanisms such as Bag-of-Words and n-grams outperform traditional machine learning approaches?
- Will NSC Algorithm together with feature extraction mechanisms such as Bag-of-Words and n-grams outperform traditional machine learning approaches?

Proposed Method



Dataset Description

Raw MIME format Email files

- Initially, the dataset comprised of 500,000 files from 150 authors which included files from sent items, sent mail, all documents, contact, inbox, deleted items, etc
- These files contained emails sent by the author, sent to the author, automated script mails and even spam.
- Each file is encoded in MIME version 1.0 which had metadata such as message ID, email id of sender, receiver, date sent, content-type, cc, bcc, etc.


Data Extraction


- Only those emails that were composed by the author or the emails in which the metadata 'from' matched the authors name were selected and from sent, sent-items, sent mail folder category.
- Each of the file was parsed using email parser package in python to get the sender information as well as the mail body of the email.
- The set of author and their corresponding content was stored as a list which was finally stored in a CSV file with around 90000 rows for 'author' and 'content'.



Data Transformation

- Transform data to be loaded in to Package 'stylo'.





(b) Word freq. before removing stop words

Data Cleaning - Snowball Stopwords


- Snowball set of stopwords.

```
> length(stop_words_snowball)
[1] 175
> stop_words_snowball
[1] "i"      "me"      "my"      "myself"  "we"      "our"      "ours"      "ourselves"
[9] "you"    "your"    "yours"    "yourself" "yourselves" "he"      "him"      "his"
[17] "himself" "she"     "her"      "hers"     "herself"   "it"      "its"      "itself"
[25] "they"    "them"    "their"    "theirs"   "themselves" "what"    "which"    "who"
[33] "whom"    "this"    "that"     "these"    "those"     "am"      "is"      "are"
```

Data Cleaning - Smart Stopwords

- Smart set of stopwords.

```
> length(stop_words_smart)
[1] 571
> stop_words_smart
[1] "a"           "a's"         "able"        "about"       "above"       "according"
[7] "accordingly" "across"      "actually"    "after"       "afterwards"  "again"
[13] "against"     "ain't"       "all"         "allow"       "allows"      "almost"
[19] "alone"       "along"       "already"     "also"        "although"    "always"
[25] "am"          "among"       "amongst"     "an"          "and"         "another"
[31] "any"         "anybody"    "anyhow"      "anyone"      "anything"    "anyway"
```





(b) Word freq. after removing stop words


Data Sampling

- Large Dataset
- Stratified Sampling to keep distribution from this Authors.
- From 5% of Population and 10 Author Sample (≈ 1500 emails)
- Same Sampling method to get training and test set (0.75/0.25)


Data Sampling and Data Splitting




(a) Original Data




(b) 0.05% Stratified



(c) 10 Classes Stratified



(d) Training set stratified



(e) Test set stratified


Feature Extraction

Frequency Vector:

- BoW
- BoW after Snowball filter
- BoW after Smart filter
- n-gram 3000
- n-gram 6000

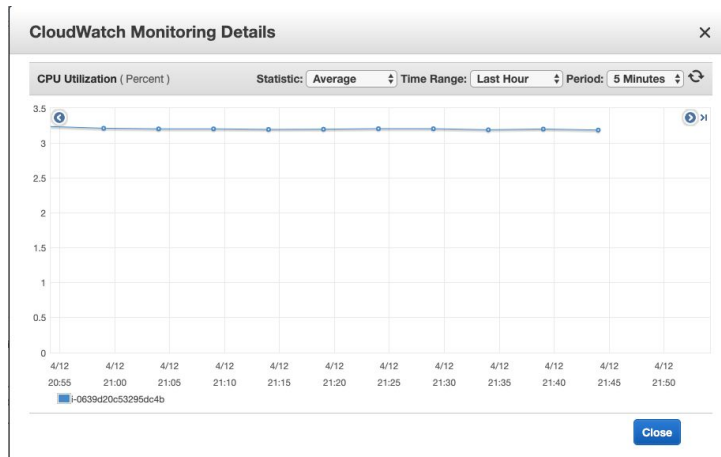
Computation-Intensive Process!!!

Feature Extraction



AWS EC2 c4.8xlarge | 36 vCPU, 60 GiB

AWS EC2 p2.xlarge | 1 GPU, 4 vCPUs, 61 GiB



Bag-of-Words (BoW)

- Representation of an observation in a tokenized way, where each word represents a feature(column).

```
> text_sample_BOW<-"This is an example of bag of words"
> print(text_sample_BOW)
[1] "This is an example of bag of words"
> tokenized_BOW<-txt.to.words.ext(text_sample_BOW, language = "English.all",
+                               preserve.case = FALSE)
> BOW<-txt.to.features(tokenized.text = tokenized_BOW,ngram.size = 1,features = 'w')
> print(BOW)
[1] "this"    "is"      "an"      "example" "of"      "bag"     "of"      "words"
```

n-gram

- Representation of an observation in a tokenized way, where each sequence of "n" words or characters represents a feature(column).

```
> text_sample_ngram<-"This is an example of n-gram"
> print(text_sample_ngram)
[1] "This is an example of n-gram"
> tokenized_ngram<-txt.to.words.ext(text_sample_ngram, language = "English.all",
+                                   preserve.case = FALSE)
> ngram<-txt.to.features(tokenized.text = tokenized_ngram,ngram.size = 3,features = 'w')
> print(ngram)
[1] "this is an"      "is an example"  "an example of"  "example of n-gram"
```

Feature Vector

- Representation of an observation as a frequency vector where each column represent a feature and its value the respective frequency of that feature.

```
> fec_vec_training_set_bow[1:3,1:10]
```

	enron	ect	com	hou	thanks	subject	cc	please	know	can
108_1038.txt	0	8	0	4	0.000000	2	2	0.000000	0.000000	0.000000
108_1086.txt	0	0	0	0	50.000000	0	0	0.000000	0.000000	0.000000
108_112.txt	0	0	0	0	2.272727	0	0	2.272727	2.272727	2.272727

Delta Algorithm and Metrics

- Variation of k-nn algorithms with different distance metrics.
- Algorithms that work on normalized vectors (rows).
- Also generate new distance metrics that normalize values over each feature(columns).
- Different types: Delta, Argamon, Eder, Eder Simple.

$$\Delta_{(AB)} = \frac{1}{n} \sum_{i=1}^n \left| \frac{A_i - \mu_i}{\sigma_i} - \frac{B_i - \mu_i}{\sigma_i} \right|$$

J. Burrows. 'Delta': a Measure of Stylistic Difference and a Guide to Likely Authorship. *Literary and Linguistic Computing*, 17(3):267–287, 9 2002.

Bug Fixed

Original

```

165 for(h in 1:length(selected.dist[,1])) {
166     ranked.c = order(selected.dist[h,])[1:no.of.candidates]
167     current.sample = classes.training.set[ranked.c[1]]
168     classification.results = c(classification.results, current.sample)
169     #
170     current.ranking = classes.training.set[ranked.c]
171     current.scores = selected.dist[h,ranked.c]
172     classification.scores = rbind(classification.scores, current.scores)
173     classification.rankings = rbind(classification.rankings, current.ranking)
174 }

```

Fixed

```

165 for(h in 1:length(selected.dist[,1])) {
166     ranked.c = order(selected.dist[h,])[1:no.of.candidates]
167
168     #current.sample = classes.training.set[ranked.c[1]]
169     #classification.results = c(classification.results, current.sample)
170     #
171     current.ranking = classes.training.set[ranked.c]
172
173     #Begin P09 JuanJoseRivera Modification
174     freq_table<-table(current.ranking)
175     current.sample<-names(freq_table)[which (freq_table==max(freq_table))][1]
176     classification.results = c(classification.results, current.sample)
177     #End P09 JuanJoseRivera Modification
178
179     current.scores = selected.dist[h,ranked.c]
180     classification.scores = rbind(classification.scores, current.scores)
181     classification.rankings = rbind(classification.rankings, current.ranking)
182 }

```


Package: stylo, Function: perform.delta(), File: perform.delta.R

Nearest Shrunk Centroids (NSC)


- Variation of Nearest Centroid Classifier.
- Used in cancer diagnosis with gene expressions (high-dimensional representation)
- High accuracy in that domain.
- Feature Vector could be considered a high-dimensional space.

R. Tibshirani, T. Hastie, B. Narasimhan, and G. Chu. Diagnosis of multiple cancer types by shrunken centroids of gene expression. Technical report, 2002


Classification Comparison (Accuracy)




(a) w/ stop words




(b) w/o Snowball



(c) w/o Smart



(d) n-gram 3000



(e) n-gram 6000

Classification Results (Accuracy)

Algorithm	Accuracy				
	BoW			n-gram	
	BoW	Snowball	Smart	3000	6000
Delta-based Algorithms					
Delta.Delta	0.1942257	0.2204724	0.2729659	0.3494624	0.311828
Delta.Eder	0.2020997	0.2152231	0.2808399	0.3709677	0.3225806
Delta.Argamon	0.2493438	0.2388451	0.3175853	0.3602151	0.3225806
Delta.Simple	0.3569554	0.3175853	0.3700787	0.4301075	0.4112903
Delta.Cosine	0.5879265	0.7007874	0.7270341	0.5510753	0.5430108
Delta.Wurzburg	0.5931759	0.5748031	0.5800525	0.5376344	0.4892473
Delta.Entropy	0.335958	0.2965879	0.3595801	0.422043	0.3978495
Delta.Euclidean*	0.4934383	0.488189	0.5301837	0.4677419	0.4301075
Delta.Manhattan*	0.5380577	0.2834646	0.328084	0.4005376	0.3844086
Delta.Canberra*	0.6456693	0.7034121	0.7217848	0.5698925	0.5537634
Non Delta-based Algorithms					
K-NN	0.4908136	0.5013123	0.5459318	0.4892473	0.4301075
SVM	0.5984252	0.5616798	0.5564304	0.4489247	0.3575269
NSC	0.7454068	0.7506562	0.7427822	0.5349462	0.483871
Naive-Bayes	0.1181102	0.1128609	0.1049869	0.07795699	0.08333333

Conclusions

- **NSC** with BoW performed best.
- **Delta.Canberra** with BoW is the next best.
- For BoW the removal of stopwords, improves the accuracy of the classifiers that shows high accuracy(Delta.Cosine, Delta.Canberra, NSC).
- According to theory, Delta Algorithms should perform better than the results of our experiments.
- We think that the improvement of our feature selection techniques and feature engineering could improve the current experiment results for Delta and NSC classifiers.
- Do not underestimate data preprocessing and transformation!!!

Future Improvement Plans

- Feature Selection
- Feature Engineering
- NLP-based Features

Questions????